

# Computer Architecture - Fall 2016

## Lab 1

### Introduction

The purpose of this lab is to reinforce the concepts of assembly language and assemblers. In this lab assignment, you will write an LC-3b Assembler, whose job will be to translate assembly language source code into the machine language (ISA) of the LC-3b. You will also write a program to solve a problem in the LC-3b Assembly Language.

In Lab Assignments 2 and 3, you will close the loop by completing the design of two types of simulators for the LC-3b, and test your assembler by having the simulators execute the program you wrote and assembled in this lab.

### Part I: Write an assembler for the LC-3b Assembly Language

The general format of a line of assembly code, which will be the input to your assembler, is as follows:

```
label opcode operands ; comments
```

The leftmost field on a line will be the label field. Valid labels consist of one to twenty alphanumeric characters (i.e., a letter of the alphabet, or a decimal digit), starting with a letter of the alphabet. A valid label cannot be the same as an opcode or a pseudo-op. A valid label must start with a letter other than 'x' and consist solely of alphanumeric characters – a to z, 0 to 9. The label is optional, i.e., an assembly language instruction can leave out the label. **A valid label cannot be IN, OUT, GETC, or PUTS. The entire assembly process, including labels, is case-insensitive.** A label is necessary if the program is to branch to that instruction or if the location contains data that is to be addressed explicitly.

The opcode field can be any one of the following instructions:

```
ADD, AND, BR(all 8 variations), HALT, JMP, JSR, JSRR, LDB, LDW,  
LEA, NOP,  
NOT, RET, LSHF, RSHFL, RSHFA, RTI, STB, STW, TRAP, XOR
```

The number of operands depends on the operation being performed. It can consist of register names, labels, or constants (immediates). If a hexadecimal constant is used, it must be prefixed with the 'x' character. Similarly, decimal constants must be prefixed with a '#' character.

Optionally, an instruction can be commented, which is good style if the comment contains

meaningful information. Comments follow the semicolon and are not interpreted by the Assembler. Note that the semicolon prefaces the comment, and a newline ends the comment. Other delimiters are not allowed.

In this lab assignment, the NOP instruction translates into the machine language instruction 0x0000.

Note that you should also implement the HALT instruction as TRAP x25. Other TRAP commands (GETC, IN, OUT, PUTS) need not be recognized by your assembler for this assignment.

In addition to LC-3b instructions, an assembly language also contains pseudo-ops, sometimes called macro directives. These are messages from the programmer to the assembler that assist the assembler in performing the translation process. In the case of our LC-3b Assembly Language, we will only require three pseudo-ops to make our lives easier: .ORIG, .END, and .FILL.

An assembly language program will consist of some number of assembly language instructions, delimited by .ORIG and .END. The pseudo-op .END is a message to the assembler designating the end of the assembly language source program. The .ORIG pseudo-op provides two functions: it designates the start of the source program, and it specifies the location of the first instruction in the object module to be produced. For example, .ORIG N means “the next instruction will be assigned to location N.” The pseudo-op .FILL W assigns the value W to the corresponding location in the object module. W is regarded as a word (16-bit value) by the .FILL pseudo-op.

The task of the assembler is that of line-by-line translation. The input is an assembly language file, and the output is an object (ISA) file (consisting of hexadecimal digits). To make it a little more concrete, here is a sample assembly language program:

```
;This program counts from 10 to 0
    .ORIG x3000
    LEA R0, TEN           ;This instruction will be loaded into memory location x3000
    LDW R1, R0, #0
START ADD R1, R1, #-1
    BRZ DONE
    BR START

                                ;blank line
DONE  TRAP x25              ;The last executable instruction
TEN   .FILL x000A           ;This is 10 in 2's comp, hexadecimal
```

```
.END                ;The pseudo-op, delimiting the source program
```

And its corresponding ISA program:

```
0x3000
0xE005
0x6200
0x127F
0x0401
0x0FFD
0xF025
0x000A
```

Note that each line of the output is a four digit hex number, prefixed with “0x”, representing the 16-bit machine instruction. The reason that your output should be prefixed with “0x” is because the simulator for Lab Assignment 2 that you will write in C expects the input data to be expressed in hex, and C syntax requires hex data to start with "0x". Also note that BR instruction is assembled as the unconditional branch, BRnzp.

When this program is loaded into the simulator, the instruction 0xE005 will be loaded into the memory location specified by the first line of the program, which is x3000. As instructions consist of two bytes, the second instruction, 0x6200, will be loaded into memory location x3002. Thus, memory locations x3000 to x300C will contain the program.

We have included below another example of an assembly language program, and the result of the assembly process. In this case, the .ORIG pseudo-op tells the assembler to place the program at memory address #4096.

```

        .ORIG #4096
A       LEA R1, Y
        LDW R1, R1, #0
        LDW R1, R1, #0
        ADD R1, R1, R1
        ADD R1, R1, x-10    ;x-10 is the negative of x10
        BRN A
        HALT
Y       .FILL #263
        .FILL #13
        .FILL #6
        .END
```

would be assembled into the following:

```
0x1000
0xE206
0x6240
0x6240
0x1241
0x1270
0x09FA
0xF025
0x0107
0x000D
0x0006
```

Important note: even though this program will assemble correctly, it may not do anything useful.

## The Assembly Process

Your assembler should make two passes of the input file. In the first pass, all the labels should be bound to specific memory addresses. You create a symbol table to contain those bindings. Whenever a new instruction label is encountered in the input file, it is assigned to the current memory address.

The second pass performs the translation from assembly language to machine language, one line at a time. It is during this pass that the output file should be generated.

You should write your program to take two command-line arguments. The first argument is the name of a file that contains a program written in LC-3b assembly language, which will be the input to your program. The second argument is the name of the file to which your program will write its output. In other words, this is the name of the file which will contain the LC-3b machine code corresponding to the input assembly language file. For example, we should be able to run your assembler with the following command-line input:

```
assemble <source.asm> <output.obj>
```

where **assemble** is the name of the executable file corresponding to your compiled and linked program; **source.asm** is the input assembly language file, and **output.obj**; is the output file that will contain the assembled code.

You will need to include some basic error checking within your assembler to handle improperly constructed assembly language programs. Your assembler must detect three types of errors and must return three different error codes. The errors to be detected are *undefined labels* (error code 1), *invalid opcodes* (error code 2), and *invalid constants* (error code 3). An invalid constant is a constant that is too large to be assembled into an LC-3b instruction. If the `.ORIG` pseudo-op contains an address that is greater than an address that can be represented in the 16-bit address space, your program should return error code 3. Also, if the

`.ORIG` statement specifies an address that is not word-aligned, your program should return error code 3. Your program must return the error codes via the `exit(n)` function, where `n` denotes the error code number. If the assembly language program does not contain any errors, you must exit with error code 0. Exiting with the correct codes is very important since they will be used in the grading process. On Linux, you can determine the exit code of your assembler by executing `echo $?` right after running the assembler.

This error checking is the bare minimum that we expect. You can return error code 4 for any other errors you find. Just be sure that the errors don't fall within the first three categories specified above.

## Examples of error codes

### Error code 1: undefined label

A label is used by an instruction but the label is not in the symbol table, e.g.

- ```
.ORIG x3000
    LEA R0, DATA1    ; DATA1 is not defined in the assembly code
.END
```
- ```
.ORIG x3000
    JSR ADD           ; JSR is parsed as an opcode, and ADD is an
                        ; illegal label. While illegal labels
                        ; should return error code 4, we accept
                        ; error code 1 for this case, too.
.END
```

### Error code 2: invalid opcode

An invalid opcode is one that is not defined in the LC-3b ISA, e.g.

- ```
.ORIG x1000
    MUL R0, R1, R2
.END
```
- ```
.ORIG x1000
    ABC
.END
```

### Error code 3: invalid constant

An invalid constant is a constant that is too large to be assembled into an LC-3b instruction. An odd constant that follows `.ORIG` is also an invalid constant.

- ```
.ORIG x1000
    ADD R0, R1, #20    ; error
.END
```
- ```
.ORIG x1001           ; error
    ADD R0, R1, #1
.END
```

### Error code 4: other error

These errors which do not belong to any of the above categories.

Examples:

- ```
.ORIG x1000
ADD R0, R1      ; wrong number of operands
.END
```
- ```
.ORIG x1000
.FILL           ; missing operand
.END
```
- ```
.ORIG x1000
ADD R1, #2, R3   ; unexpected operand
.END
```
- ```
.ORIG x1000
ADD R9, R0, #1    ; R9 is an invalid register number
.END
```
- ```
.ORIG x1000
ADD R1, R0, 1; 1 is an invalid operand (neither a register nor an immediate)
.END
```

**If a label and an instruction that uses it are too far apart and the offset cannot be specified properly in the machine code, you should produce error code 4.**

Your assembler should accept an “empty” program, i.e. one with just a valid `.ORIG` and a `.END`. E.g. the following assembly program would be assembled to only one line containing the starting address (`0x3000`).

```
.ORIG x3000
.END
```

Note: your assembler needs to recognize only labels as operands for `LEA`, `BR`, and `JSR` instructions. For example, if the following line is in an input assembly language program, your assembler should exit with error code 4:

```
LEA R1, x100
```

## Part II: Write a **procedure** to solve the following problem

As we discussed in class, the LC-3b does not support unaligned memory accesses. Your assignment: Write a **procedure** that will load a 16 bit integer from memory into R0, regardless whether the address is aligned or unaligned. If the address is aligned, it will be done with a single load word instruction. If the address is unaligned, it will be done with two load byte instructions.

R1 contains the address of the memory location containing the low 8 bits of data to be loaded into R0. Execution should start at location x4000.

You will have no way to determine if your assembly language code works (yet!), but you can use it to determine if your assembler works. Despite this, **we will grade this program for correctness.**

**Note:** We asked you to write a procedure, not a program. Procedures terminate with a “RET” instruction, not a “HALT”. Please terminate your program with a “RET”.

## Requirements

**Important note:** because we will be evaluating your code in Unix, please be sure your code compiles using gcc with the `-ansi` flag. This means that you need to write your code in C such that it conforms to the ANSI C standard.

You can use the following command to compile your code:

```
gcc -ansi -o assemble assembler.c
```

You should also make sure that your code runs correctly on one of the ECE linux machines.

To complete Lab Assignment 1, you will need to turn in the following:

1. A C file called "assembler.c" containing an adequately documented listing of your LC-3b Assembler.
2. A source listing (LC-3b Assembly Language) of the program described above called "loadUnaligned.asm".

Submission instructions are posted here.

## Things to watch for:

Be sure that your assembler can handle comments on any line, including lines that contain pseudo-ops and lines that contain only comments. Be careful with comments that follow a HALT, NOP or RET instructions – these instructions take no operand.

Your assembler should allow hexadecimal and decimal constants after both ISA instructions, like ADD, and pseudo-ops, like .FILL.

The whole assembly process is case insensitive. That is, the labels, opcodes, operands, and pseudo-ops can be in upper case, lower case, or both, and are still interpreted the same. The parser function given in the useful code page converts every line into lower case before parsing

it.

You can assume that there will be at most 255 labels in an assembly program. You can also assume that the number of characters on a line will not exceed 255.

Your assembler needs to support all 8 variations of BR:

|            |             |
|------------|-------------|
| BRn LABEL  | BRz LABEL   |
| BRp LABEL  | BRnz LABEL  |
| BRnp LABEL | BRzp LABEL  |
| BR LABEL   | BRnzp LABEL |

## Lab Assignment 1 Clarifications

**NOTE: FAQ's for this semester will be posted here. Please check back regularly.**

1. Constants can be expressed in hex or in decimal. Hex constants consist of an 'x' or 'X' followed by one or more hex digits. Decimal constants consist of a '#' followed by one or more decimal digits. Negative constants are identified by a minus sign immediately after the 'x' or '#'. For example, #-10 is the negative of decimal 10 (i.e., -10), and x-10 is the negative of x10 (i.e. -16).
2. Since the sign is explicitly specified, the rest of the constant is treated as an unsigned number. For example, x-FF is equivalent to -255. The 'x' tells us the number is in hex, the '-' tells us it is a negative number, and "FF" is treated as an unsigned hex number (i.e., 255). Putting it all together gives us -255.
3. Your assembler does not have to check for multiple .ORIG pseudo-ops.
4. Since the .END pseudo-op is used to designate the end of the assembly language file, your assembler does not need to process anything that comes after the .END.
5. The trap vector for a TRAP instruction and the shift amount for SHF instructions must be non-negative values. If they are not, you should return error code 3.
6. The same label should not appear in the symbol table more than once. During pass 1 of the assembly process, you should check to make sure a label is not already in the symbol table before adding it to the symbol table. If the label is already in the symbol table, you should return error code 4.
7. An invalid label (i.e., one that contains non-alphanumeric characters, or one that starts



with the letter 'x' or a number) is another example of error code 4.

8. The standard C function [isalnum\(\)](#) can be used to check if a character is alphanumeric.
9. After you have gone through the input file for pass 1 of the assembler and your file pointer is at the end of the file, there are two ways you can get the file pointer back to the beginning. You can either close and reopen the file or you can use the standard C I/O function [rewind\(\)](#).
10. The following definitions can be used to create your symbol table:

```
#define MAX_LABEL_LEN 20
#define MAX_SYMBOLS 255
typedef struct {
    int address;
    char label[MAX_LABEL_LEN + 1]; /* Question for the reader: Why do we
need to add 1? */
} TableEntry;
TableEntry symbolTable[MAX_SYMBOLS];
```

11. To check if two strings are the same, you can use the standard C string function [strcmp\(\)](#). To copy one string to another, you can use the standard C string function [strcpy\(\)](#).
12. If you decide to use any of the math functions in math.h, you also have to link the math library by using the command:

```
gcc -lm -ansi -o assemble assembler.c
```

13. When your assembler finds an error in the input assembly language program, it is not required that you print out an error message to the screen. If you choose to do this to make debugging easier, that is fine. What is required is that you exit with the appropriate error code. This is what we will be checking for when we grade your program; we will ignore anything that is printed to the screen.
14. An assembly program which starts with comments before `.ORIG` is valid and your assembler should ignore them. You can assume that there will be no label in front of `.ORIG` and `.END` in the same line.
15. Your assembler needs to be able to assemble programs which begin at any point in the

LC-3b's 16-bit address space. While user programs start from x3000 and continue until xFDFF, the assembler could be used to assemble system code as well. The assembler doesn't have enough information when it is assembling the program to determine how the program will be used. In future labs, we will develop what happens if a user tries to access a protected region of memory.

16. .FILL can take an address, signed number, or unsigned number.
17. The trap vector for a TRAP instruction should be a hex number.
18. If a single line in the program contains multiple errors, feel free to exit with the appropriate error code for any of the errors that line the program contains.